

Project III: Kernel PCA and Association Rule

Appiah Prince* University of Texas at El Paso (UTEP)

October 11, 2022

Contents

1	Bring in and examine the data	2
1.1	Bring in both the train and the test data	2
1.2	Checking for columns that are unary or close to unary	2
1.3	Checking for missing values	6
2	Ordinary Principal Components Analysis (PCA)	7
3	Kernel PCA	12
4	PCA and KPCA on the test data	20
5	ASSOCIATION RULES	24
5.1	Read in Data	24
5.2	Perform frequent itemsets and association rule analysis.	25
5.3	Top 5 rules in decreasing order of confidence (conf) for item sets of size/length 2 or 3.	30
5.4	Top 5 rules in decreasing order of the lift measure for item sets of size 2 or 3.	31
5.5	Conviction measures for the top-lift 5 rules in Part (d)	31

*pappiah@miners.utep.edu

1 Bring in and examine the data

1.1 Bring in both the train and the test data

```
# BRING IN THE DATA
train <- read.table(file=
"http://archive.ics.uci.edu/ml/machine-learning-databases/optdigits/optdigits.tra",
sep=",", header = FALSE, na.strings = c("NA", "", " "),
col.names = c(paste("x", 1:64, sep=""), "digit"))
test <- read.table(file=
"http://archive.ics.uci.edu/ml/machine-learning-databases/optdigits/optdigits.tes",
sep=",", header = FALSE, na.strings = c("NA", "", " "),
col.names = c(paste("x", 1:64, sep=""), "digit"))

dim(train)
```

```
## [1] 3823 65
```

```
dim(test)
```

```
## [1] 1797 65
```

1.2 Checking for columns that are unary or close to unary

```
library(caret)
nearZeroVar(train[, -65], uniqueCut = 10, saveMetrics = TRUE)
```

```
##      freqRatio percentUnique zeroVar  nzv
## x1      0.000000      0.02615747   TRUE  TRUE
## x2     10.912162      0.23541721  FALSE FALSE
## x3      3.096296      0.44467696  FALSE FALSE
## x4      1.591376      0.44467696  FALSE FALSE
## x5      1.742919      0.44467696  FALSE FALSE
## x6      3.288235      0.44467696  FALSE FALSE
## x7     14.912821      0.44467696  FALSE FALSE
## x8    132.464286      0.41851949  FALSE  TRUE
## x9   3820.000000      0.10462987  FALSE  TRUE
## x10      7.327703      0.41851949  FALSE FALSE
## x11      1.950324      0.44467696  FALSE FALSE
## x12      2.312020      0.44467696  FALSE FALSE
```

## x13	2.734628	0.44467696	FALSE	FALSE
## x14	1.256849	0.44467696	FALSE	FALSE
## x15	11.789474	0.44467696	FALSE	FALSE
## x16	108.264706	0.39236202	FALSE	TRUE
## x17	953.500000	0.13078734	FALSE	TRUE
## x18	6.856089	0.44467696	FALSE	FALSE
## x19	1.407407	0.44467696	FALSE	FALSE
## x20	1.744000	0.44467696	FALSE	FALSE
## x21	1.458861	0.44467696	FALSE	FALSE
## x22	1.679339	0.44467696	FALSE	FALSE
## x23	10.776256	0.44467696	FALSE	FALSE
## x24	187.800000	0.23541721	FALSE	TRUE
## x25	954.750000	0.05231494	FALSE	TRUE
## x26	6.376667	0.44467696	FALSE	FALSE
## x27	1.296804	0.44467696	FALSE	FALSE
## x28	1.592129	0.44467696	FALSE	FALSE
## x29	1.497948	0.44467696	FALSE	FALSE
## x30	1.368601	0.44467696	FALSE	FALSE
## x31	10.923077	0.44467696	FALSE	FALSE
## x32	476.625000	0.07847240	FALSE	TRUE
## x33	763.600000	0.05231494	FALSE	TRUE
## x34	10.218750	0.41851949	FALSE	FALSE
## x35	1.430473	0.44467696	FALSE	FALSE
## x36	1.457386	0.44467696	FALSE	FALSE
## x37	2.160878	0.44467696	FALSE	FALSE
## x38	1.200000	0.44467696	FALSE	FALSE
## x39	5.162319	0.39236202	FALSE	FALSE
## x40	0.000000	0.02615747	TRUE	TRUE
## x41	344.000000	0.20925974	FALSE	TRUE
## x42	11.033473	0.44467696	FALSE	FALSE
## x43	2.636861	0.44467696	FALSE	FALSE
## x44	1.878834	0.44467696	FALSE	FALSE
## x45	1.222982	0.44467696	FALSE	FALSE
## x46	1.223035	0.44467696	FALSE	FALSE
## x47	7.383333	0.44467696	FALSE	FALSE
## x48	117.968750	0.13078734	FALSE	TRUE
## x49	189.550000	0.20925974	FALSE	TRUE
## x50	10.383513	0.44467696	FALSE	FALSE
## x51	1.899743	0.44467696	FALSE	FALSE
## x52	3.366534	0.44467696	FALSE	FALSE
## x53	3.143396	0.44467696	FALSE	FALSE
## x54	1.010249	0.44467696	FALSE	FALSE
## x55	9.602094	0.44467696	FALSE	FALSE
## x56	58.370968	0.28773215	FALSE	TRUE
## x57	3822.000000	0.05231494	FALSE	TRUE

```
## x58 17.326425 0.28773215 FALSE FALSE
## x59 3.471774 0.44467696 FALSE FALSE
## x60 1.907950 0.44467696 FALSE FALSE
## x61 2.320930 0.44467696 FALSE FALSE
## x62 2.797222 0.44467696 FALSE FALSE
## x63 9.741313 0.44467696 FALSE FALSE
## x64 53.970149 0.41851949 FALSE TRUE
```

```
nearZeroVar(train[, -65], names = TRUE)
```

```
## [1] "x1" "x8" "x9" "x16" "x17" "x24" "x25" "x32" "x33" "x40" "x41" "x48"
## [13] "x49" "x56" "x57" "x64"
```

```
library(caret)
nearZeroVar(test[, -65], uniqueCut = 10, saveMetrics = TRUE)
```

```
##      freqRatio percentUnique zeroVar  nzv
## x1      0.000000      0.0556483   TRUE  TRUE
## x2     11.960938      0.5008347  FALSE FALSE
## x3      3.006993      0.9460211  FALSE FALSE
## x4      1.735160      0.9460211  FALSE FALSE
## x5      1.800000      0.9460211  FALSE FALSE
## x6      2.987879      0.9460211  FALSE FALSE
## x7     15.211111      0.9460211  FALSE FALSE
## x8    145.750000      0.8903728  FALSE  TRUE
## x9    447.500000      0.1669449  FALSE  TRUE
## x10     8.540323      0.9460211  FALSE FALSE
## x11     2.192308      0.9460211  FALSE FALSE
## x12     2.193237      0.9460211  FALSE FALSE
## x13     2.575540      0.9460211  FALSE FALSE
## x14     1.361011      0.9460211  FALSE FALSE
## x15    10.067797      0.9460211  FALSE FALSE
## x16   109.250000      0.7234279  FALSE  TRUE
## x17   597.666667      0.1669449  FALSE  TRUE
## x18      6.770992      0.9460211  FALSE FALSE
## x19     1.675556      0.9460211  FALSE FALSE
## x20     1.483193      0.9460211  FALSE FALSE
## x21     1.513605      0.9460211  FALSE FALSE
## x22     2.016529      0.9460211  FALSE FALSE
## x23    10.924528      0.9460211  FALSE FALSE
## x24   147.000000      0.4451864  FALSE  TRUE
## x25   897.500000      0.1112966  FALSE  TRUE
## x26      6.287770      0.8903728  FALSE FALSE
```

```
## x27    1.244648    0.9460211  FALSE FALSE
## x28    1.388462    0.9460211  FALSE FALSE
## x29    1.549521    0.9460211  FALSE FALSE
## x30    1.504202    0.9460211  FALSE FALSE
## x31     9.705357    0.8903728  FALSE FALSE
## x32  448.250000    0.1112966  FALSE  TRUE
## x33     0.000000    0.0556483    TRUE  TRUE
## x34   10.915789    0.8347245  FALSE FALSE
## x35     1.354167    0.9460211  FALSE FALSE
## x36     1.352941    0.9460211  FALSE FALSE
## x37     1.894545    0.9460211  FALSE FALSE
## x38     1.070470    0.9460211  FALSE FALSE
## x39     4.807910    0.8347245  FALSE FALSE
## x40     0.000000    0.0556483    TRUE  TRUE
## x41  357.600000    0.2782415  FALSE  TRUE
## x42     8.862595    0.9460211  FALSE FALSE
## x43     2.284698    0.9460211  FALSE FALSE
## x44     1.823529    0.9460211  FALSE FALSE
## x45     1.300912    0.9460211  FALSE FALSE
## x46     1.698020    0.9460211  FALSE FALSE
## x47     8.714286    0.9460211  FALSE FALSE
## x48  147.916667    0.3895381  FALSE  TRUE
## x49  896.500000    0.2225932  FALSE  TRUE
## x50     9.622378    0.8347245  FALSE FALSE
## x51     2.148810    0.9460211  FALSE FALSE
## x52     2.562500    0.9460211  FALSE FALSE
## x53     2.645161    0.9460211  FALSE FALSE
## x54     1.283217    0.9460211  FALSE FALSE
## x55     9.559140    0.9460211  FALSE FALSE
## x56    52.531250    0.6677796  FALSE  TRUE
## x57 1796.000000    0.1112966  FALSE  TRUE
## x58    17.533333    0.5564830  FALSE FALSE
## x59     3.162963    0.9460211  FALSE FALSE
## x60     2.365854    0.9460211  FALSE FALSE
## x61     2.724138    0.9460211  FALSE FALSE
## x62     2.590426    0.9460211  FALSE FALSE
## x63     8.822222    0.9460211  FALSE FALSE
## x64    58.172414    0.9460211  FALSE  TRUE
```

```
nearZeroVar(test[, -65], names = TRUE)
```

```
## [1] "x1" "x8" "x9" "x16" "x17" "x24" "x25" "x32" "x33" "x40" "x41" "x48"
## [13] "x49" "x56" "x57" "x64"
```

```
train <- train[,-c(1,8,9,16,17,24,25,32,33,40,41,48,49,56,57,64)]  
test  <- test[,-c(1,8,9,16,17,24,25,32,33,40,41,48,49,56,57,64)]
```

Remarks

I removed the columns 1,8,9,16,17,24,25,32,33,40,41,48,49,56,57,64 from the columns of both train and test data since they have some values that are unary and some close to unary.

1.3 Checking for missing values

```
library(questionr)  
freq.na(train)
```

```
##      missing %  
## x2          0 0  
## x3          0 0  
## x4          0 0  
## x5          0 0  
## x6          0 0  
## x7          0 0  
## x10         0 0  
## x11         0 0  
## x12         0 0  
## x13         0 0  
## x14         0 0  
## x15         0 0  
## x18         0 0  
## x19         0 0  
## x20         0 0  
## x21         0 0  
## x22         0 0  
## x23         0 0  
## x26         0 0  
## x27         0 0  
## x28         0 0  
## x29         0 0  
## x30         0 0  
## x31         0 0  
## x34         0 0  
## x35         0 0
```


```
## x36      0 0
## x37      0 0
## x38      0 0
## x39      0 0
## x42      0 0
## x43      0 0
## x44      0 0
## x45      0 0
## x46      0 0
## x47      0 0
## x50      0 0
## x51      0 0
## x52      0 0
## x53      0 0
## x54      0 0
## x55      0 0
## x58      0 0
## x59      0 0
## x60      0 0
## x61      0 0
## x62      0 0
## x63      0 0
## digit    0 0
```

Remarks

There are no missing values in the train data.


2 Ordinary Principal Components Analysis (PCA)

```
# Parallel Boxplot of the attributes of the train data
boxplot(train[, -49], col = rainbow(ncol(train[, -49])), main="Boxplot of train data")
```

Boxplot of train data

```
# Parallel Boxplot of the attributes of the test data
```

```
boxplot(test[, -49], col = rainbow(ncol(test[, -49])), main="Boxplot of test data")
```

Boxplot of test data

Remarks

- Majority of the predictors in both the train and test data have unequal range and unequal variation.
- Hence, scaling is necessary for some modeling approaches.

```
# scaling the train and test data
```

```
train_scaled <- data.frame(apply(train[, -49], 2, scale, center=T, scale=T))
```


```
mean <- apply(train_scaled, 2, mean)
```

```
sd <- apply(train_scaled, 2, sd)
```

```
test_scaled <- data.frame(scale(test[, -49], center = mean, scale = sd))
```


```
boxplot(train_scaled, col = rainbow(ncol(train_scaled)), main="Boxplot of standardized
```

Boxplot of standardized train data



```
boxplot(test_scaled, col = rainbow(ncol(test_scaled)), main="Boxplot of standardized tes
```

Boxplot of standardized test data



Remarks

After scaling both the test and train data, we see that very few of the attributes of test and train data have unequal range and variation. Hence, we can now run the ordinary principal components analysis (PCA).


```
pca <- prcomp(train_scaled, retx=TRUE, center=F, scale=F)

# OBTAIN EIGENVALUES
lambda <- eigen(cov(train_scaled), only.values = T)$values
lambda

## [1] 6.91007451 6.25858444 4.64555316 3.46269437 2.68099761 2.47851068
## [7] 2.04654171 1.83311085 1.55795400 1.45344719 1.25778760 1.21108619
## [13] 1.09097620 0.92376560 0.85699940 0.77122785 0.71038203 0.62704218
## [19] 0.55046754 0.53121523 0.47656857 0.45452988 0.42244094 0.38724553
## [25] 0.37911851 0.34677924 0.32034160 0.28199851 0.26768836 0.26197379
## [31] 0.25311899 0.21488696 0.20678233 0.19538461 0.17944841 0.16851168
## [37] 0.16195405 0.15399054 0.13799300 0.13102482 0.12169469 0.11560572
## [43] 0.10320643 0.10273704 0.09181544 0.07903824 0.06715225 0.05855154


#screeplot of variance
screeplot(pca, npcs = 30, type="lines", main="Scree Plot", col = "red")
```

Scree Plot



```
# PLOT FIRST TWO PCs
par(mfrow=c(1,1), mar=rep(4,4))
plot(pca$x[,1:2], pch="", main="PC.1 and PC.2 for the train handwritten digit data")
text(pca$x[,1:2], labels=train$digit, col= train$digit)
abline(v=0, lty=2)
abline(h=0, lty=2)
```

PC.1 and PC.2 for the train handwritten digit data



Remarks

We can see from the plot graph that the first two PCs fairly successfully separate the digits. We see, for instance, that most 6s lie on the top of the plot, most 4s lie on the upper right, most 7s on the bottom right, and most 2s on the middle-top left. There are however regions of overlap.

3 Kernel PCA

```
# Using different kernel functions
library(kernlab)
```

```
##
## Attaching package: 'kernlab'

## The following object is masked from 'package:ggplot2':
##
##   alpha
```

```
kernel_pca1 <- kpca(~., data=train_scaled, kernel="rbfdot", kpar=list(sigma=0.01), features=10)
kernel_pca2 <- kpca(~., data=train_scaled, kernel="vanilladot", kpar=list(),
  features = 10)
kernel_pca3 <- kpca(~., data=train_scaled, kernel="polydot", kpar=list(degree=2),
  features=10)
kernel_pca4 <- kpca(~., data=train_scaled, kernel="laplacedot",
  kpar=list(sigma=0.01), features=10)
```

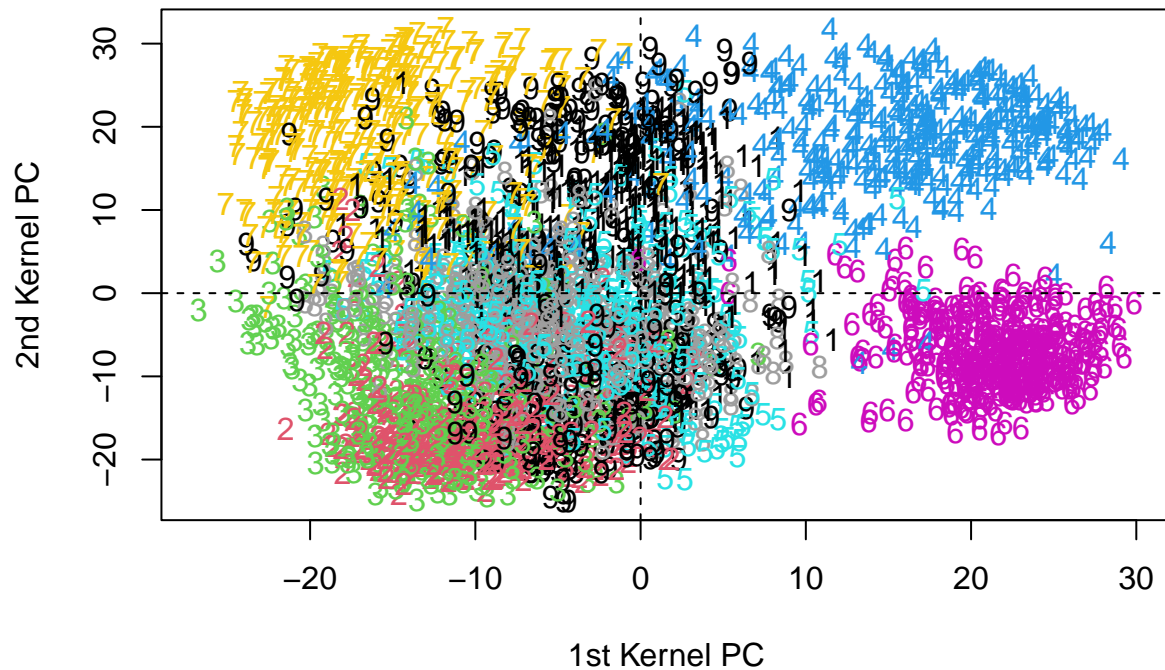
```
# Get the variance of each kernel pca.
```

```
var.pc1 <- eig(kernel_pca1)
var.pc2 <- eig(kernel_pca2)
var.pc3 <- eig(kernel_pca3)
var.pc4 <- eig(kernel_pca4)
variance <- data.frame(var.pc1, var.pc2, var.pc3, var.pc4)
variance
```

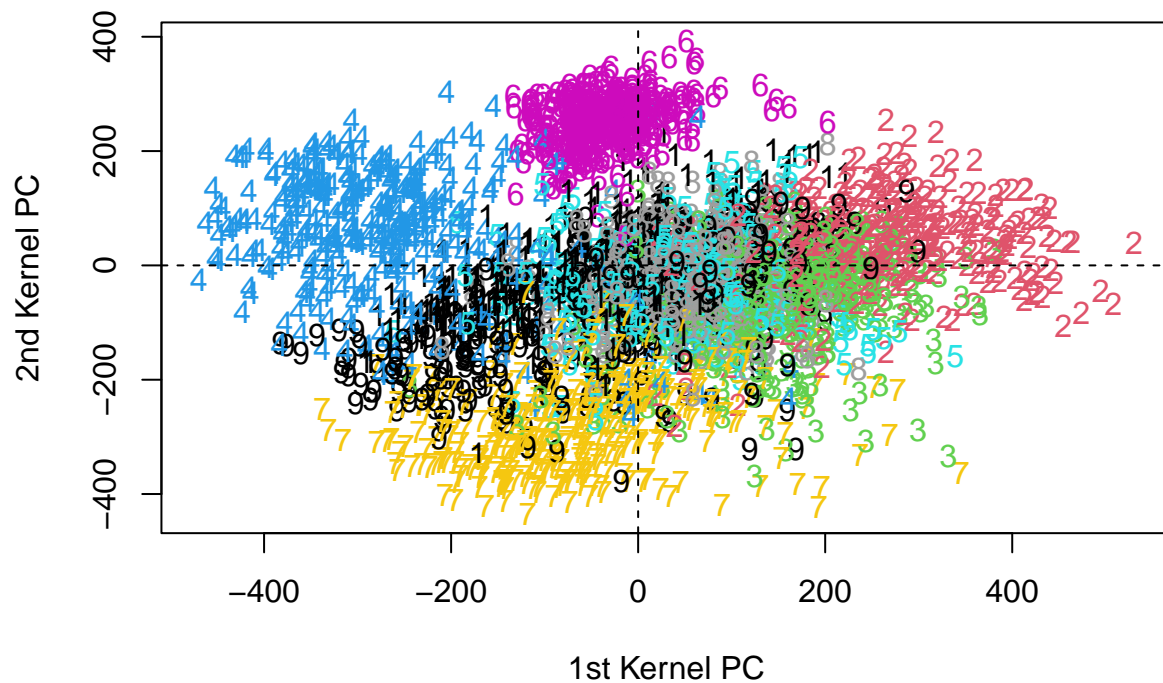
```
##           var.pc1  var.pc2  var.pc3  var.pc4
## Comp.1  0.05023886 6.908267 159.43549 0.006244509
## Comp.2  0.04913771 6.256947 139.35988 0.006058360
## Comp.3  0.04201080 4.644338 119.44195 0.004937771
## Comp.4  0.02887515 3.461789  93.99018 0.003473661
## Comp.5  0.02499205 2.680296  78.91073 0.002813834
## Comp.6  0.02282597 2.477862  75.12011 0.002615203
## Comp.7  0.02019387 2.046006  67.26845 0.002314876
## Comp.8  0.01788108 1.832631  63.87454 0.002129724
## Comp.9  0.01420419 1.557546  56.27348 0.001654528
## Comp.10 0.01301382 1.453067  43.95150 0.001544828
```

- Plotting the first two PCs for each of the kernel pca

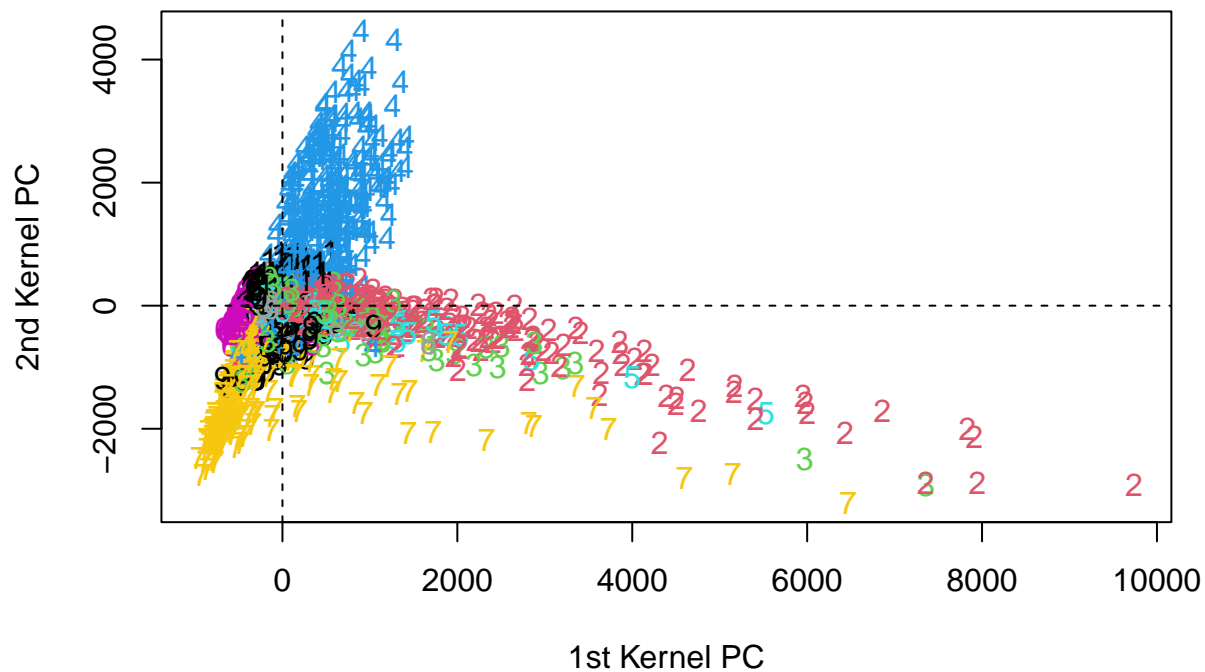
```
PC1 <- rotated(kernel_pca1) # returns the data projected in the (kernel) pca space
plot(PC1[, 1:2], col=train$digit, pch="",
  main="KPC.1 and KPC.2 for the train handwritten digit data",
  xlab="1st Kernel PC", ylab="2nd Kernel PC")
text(PC1[, 1:2], labels=train$digit, col= train$digit)
abline(v=0, lty=2)
abline(h=0, lty=2)
```

KPC.1 and KPC.2 for the train handwritten digit data


```
PC2 <- rotated(kernel_pca2)      # returns the data projected in the (kernel) pca space
plot(PC2[, 1:2], col=train$digit, pch="",
     main="KPC.1 and KPC.2 for the train handwritten digit data",
     xlab="1st Kernel PC", ylab="2nd Kernel PC")
text(PC2[, 1:2], labels=train$digit, col= train$digit)
abline(v=0, lty=2)
abline(h=0, lty=2)
```

KPC.1 and KPC.2 for the train handwritten digit data

```
PC3 <- rotated(kernel_pca3)      # returns the data projected in the (kernel) pca space
plot(PC3[, 1:2], col=train$digit, pch="",
     main="KPC.1 and KPC.2 for the train handwritten digit data",
     xlab="1st Kernel PC", ylab="2nd Kernel PC")
text(PC3[, 1:2], labels=train$digit, col= train$digit)
abline(v=0, lty=2)
abline(h=0, lty=2)
```

KPC.1 and KPC.2 for the train handwritten digit data


```
PC4 <- rotated(kernel_pca4)      # returns the data projected in the (kernel) pca space
plot(PC4[, 1:2], col=train$digit, pch="",
     main="KPC.1 and KPC.2 for the train handwritten digit data",
     xlab="1st Kernel PC", ylab="2nd Kernel PC")
text(PC4[, 1:2], labels=train$digit, col= train$digit)
abline(v=0, lty=2)
abline(h=0, lty=2)
```


KPC.1 and KPC.2 for the train handwritten digit data

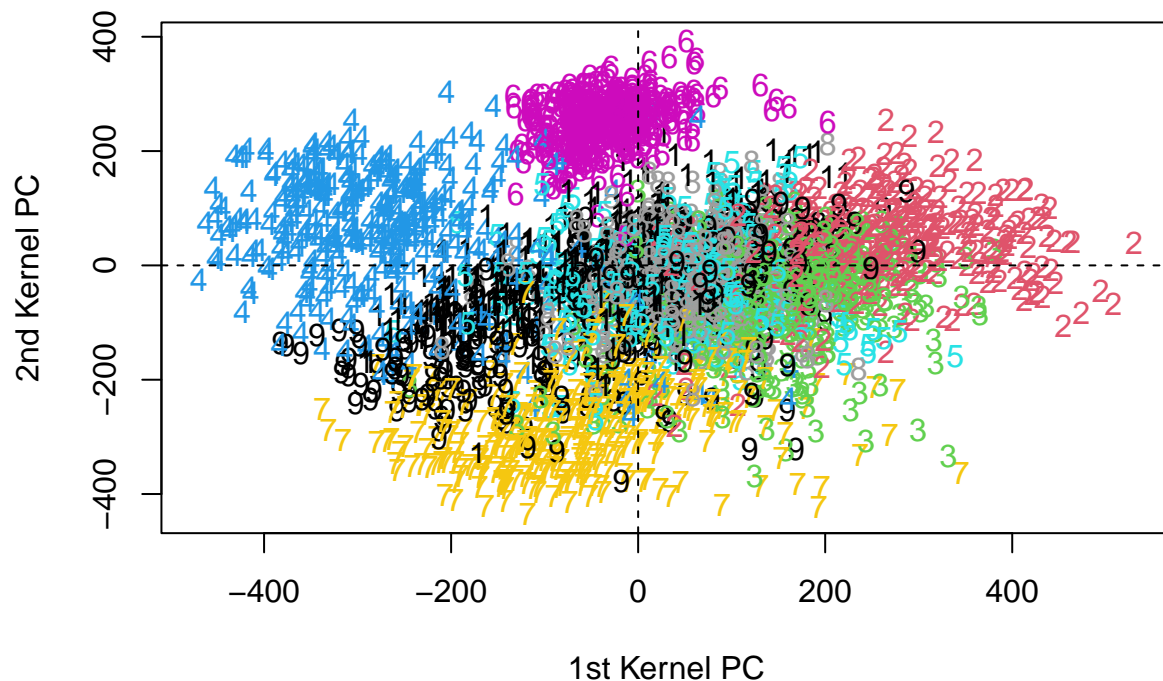
Remarks

We observed that the kernel pca using the vanilladot kernel function separated or clustered the digits well as compared to the other kernel pca's using different kernel functions. So, I choose the kernel pca using the vanilladot kernel function.

```
#screeplot for the variance of the kernel pca using vanilladot.
var.pc <- eig(kernel_pca2)
prop.pc <- var.pc/sum(var.pc)
plot(prop.pc, xlab = "Principal Component", col = "red",
      ylab = "Proportion of Variance Explained", type = "b", pch = 19)
```



```
# Plot THE DATA PROJECTION ON THE KERNEL PCS
PC <- rotated(kernel_pca2)      # returns the data projected in the (kernel) pca space
plot(PC[, 1:2], col=train$digit, pch="",
      main="KPC.1 and KPC.2 for the train handwritten digit data",
      xlab="1st Kernel PC", ylab="2nd Kernel PC")
text(PC[, 1:2], labels=train$digit, col= train$digit)
abline(v=0, lty=2)
abline(h=0, lty=2)
```

KPC.1 and KPC.2 for the train handwritten digit data

Remarks

- We can see from the above plots that the first two PCs fairly successfully separate the digits. We see, for instance, that most 6s lie on the top of the plot, most 4s on the top left, most 7s on the bottom and most 2s lie on the middle right. There are however regions of overlap.
- The choice of kernel function used is `vanilladot`(linear kernel function).
- The parameter is `degree = 1`.
- comparison of PCA and KPCA

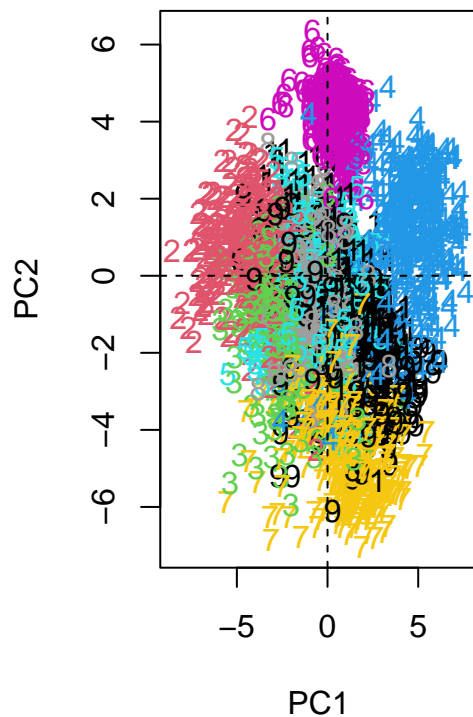

```
par(mfrow=c(1,2), mar=rep(4,4))
plot(pca$x[,1:2], pch="", main="Ordinary PCA for the train data")
text(pca$x[,1:2], labels=train$digit, col= train$digit)
abline(v=0, lty=2)
abline(h=0, lty=2)

plot(PC[, 1:2], col=train$digit, pch="",
      main="Kernel PCA for the train data",
```

```

xlab="1st Kernel PC", ylab="2nd Kernel PC")
text(PC[, 1:2], labels=train$digit, col= train$digit)
abline(v=0, lty=2)
abline(h=0, lty=2)

```

Ordinary PCA for the train data**Kernel PCA for the train data**

Remarks

We see from the above plots that there is no significant difference between clustering of the digits. Both methods show that the first two PCs explain a substantial portion of the variation in the data.

4 PCA and KPCA on the test data

```

#ordinary pca
pred_pca <- predict(pca, test_scaled)

# comparison of the PCA results on the train and test data


par(mfrow=c(1,2), mar=rep(4,4))

```


```
plot(pca$x[,1:2], pch="", main="Ordinary PCA on train data")
text(pca$x[,1:2], labels=train$digit, col= train$digit)
abline(v=0, lty=2)
abline(h=0, lty=2)

plot(pred_pca[,1:2], pch="", main="Ordinary PCA on test data")
text(pred_pca[,1:2], labels=test$digit, col= test$digit)
abline(v=0, lty=2)
abline(h=0, lty=2)
```

Ordinary PCA on train data



Ordinary PCA on test data



Remarks

There is no significant difference between ordinary pca on both the train and test data.

```
#kernel pca
pred_kernel_pca <- predict(kernel_pca2, test_scaled)

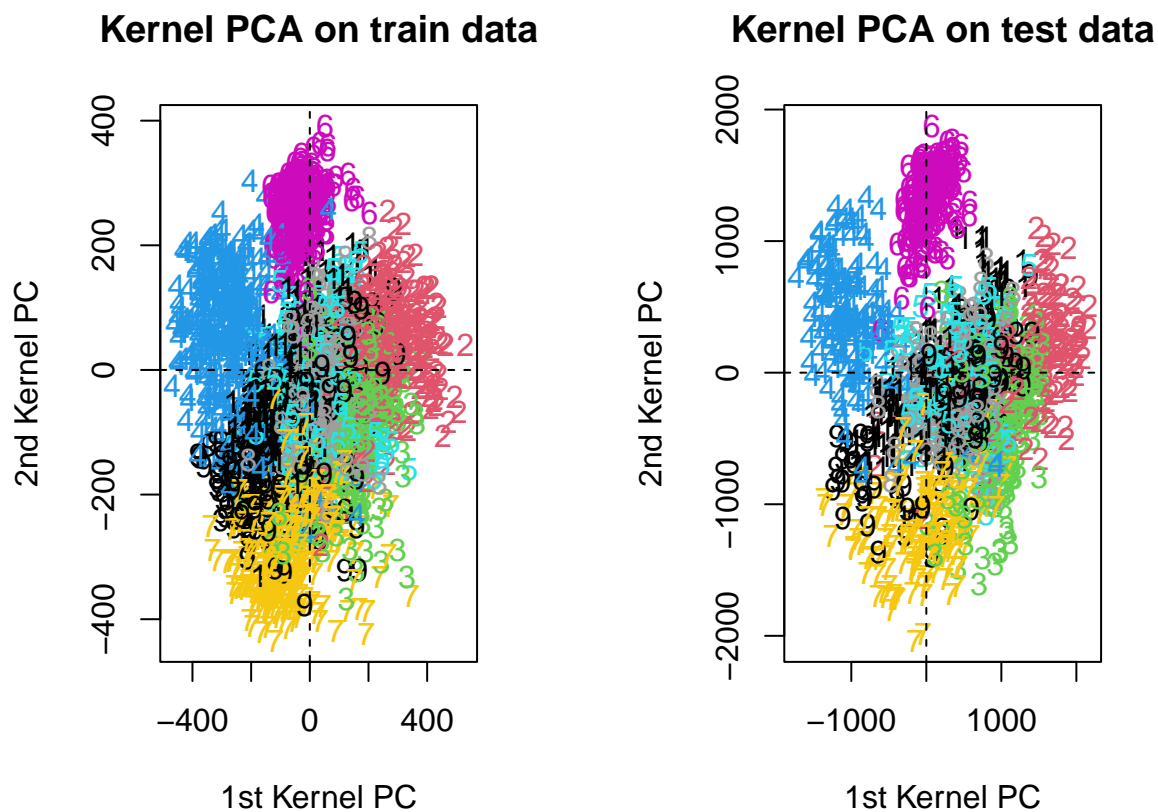
par(mfrow=c(1,2), mar=rep(4,4))
plot(PC[, 1:2], col=train$digit, pch="",
      main="Kernel PCA on train data",
```

```

    xlab="1st Kernel PC", ylab="2nd Kernel PC")
text(PC[, 1:2], labels=train$digit, col= train$digit)
abline(v=0, lty=2)
abline(h=0, lty=2)

plot(pred_kernel_pca[, 1:2], col=test$digit, pch="",
     main="Kernel PCA on test data",
     xlab="1st Kernel PC", ylab="2nd Kernel PC")
text(pred_kernel_pca[, 1:2], labels=test$digit, col= test$digit)
abline(v=0, lty=2)
abline(h=0, lty=2)

```



Remarks

There is no significant difference between the result of the Kernel pca on the train data and the kernel pca on test data.

- comparison of ordinary pca and kernel pca on the test data

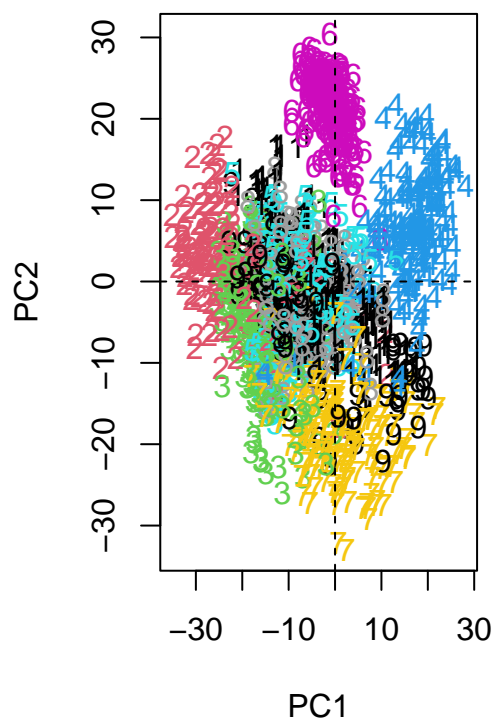

```

par(mfrow=c(1,2), mar=rep(4,4))

plot(pred_pca[,1:2], pch="", main="Ordinary PCA on test data")
text(pred_pca[,1:2], labels=test$digit, col= test$digit)
abline(v=0, lty=2)
abline(h=0, lty=2)

plot(pred_kernel_pca[, 1:2], col=test$digit, pch="",
      main="Kernel PCA on test data",
      xlab="1st Kernel PC", ylab="2nd Kernel PC")
text(pred_kernel_pca[, 1:2], labels=test$digit, col= test$digit)
abline(v=0, lty=2)
abline(h=0, lty=2)

```

Ordinary PCA on test data**Kernel PCA on test data****Remarks**

- There is no significant difference between the results of the ordinary pca and the kernel pca on the test data.
- We observe that most 2s lie on the middle left on the ordinary pca while they lie on the middle right on the kernel pca.

- In both cases they fairly separate the digits well.

5 ASSOCIATION RULES

5.1 Read in Data

```
library(arules)

## Loading required package: Matrix

##
## Attaching package: 'arules'

## The following object is masked from 'package:kernlab':
##
##      size

## The following objects are masked from 'package:base':
##
##      abbreviate, write

bible <- read.transactions(file="AV1611Bible.txt",
format = "basket", sep = " ", rm.duplicates =F,
quote="") # DOUBLE/SINGLE QUOTE ISSUE
dat <- bible; dim(dat)

## [1] 31101 12767

inspect(dat[1:5, ])
```

```
##      items
## [1] {beginning,
##      created,
##      earth,
##      god,
##      heaven}
## [2] {darkness,
##      deep,
##      earth,
```



```
##      face,
##      form,
##      god,
##      moved,
##      spirit,
##      upon,
##      void,
##      waters,
##      without}
## [3] {god,
##      let,
##      light,
##      said,
##      there}
## [4] {darkness,
##      divided,
##      god,
##      good,
##      light,
##      saw}
## [5] {called,
##      darkness,
##      day,
##      evening,
##      first,
##      god,
##      light,
##      morning,
##      night}
```


5.2 Perform frequent itemsets and association rule analysis.

```
# The first 15 items (frequency/support)
itemFrequency(dat[, 1:15])
```

```
##      aaron      aaron's    aaronites    abaddon      abagtha      abana
## 9.742452e-03 9.967525e-04 6.430661e-05 3.215331e-05 3.215331e-05 3.215331e-05
##      abarim      abase      abased      abasing      abated      abba
## 1.286132e-04 1.286132e-04 1.286132e-04 3.215331e-05 1.929198e-04 9.645992e-05
##      abda      abdeel      abdi
## 6.430661e-05 3.215331e-05 9.645992e-05
```

```
# Plot items with high frequencies.
```

```
itemFrequencyPlot(dat, topN=10, support = 0.01, cex.names = 0.8, col="blue")
```



Remarks

We observe that lord has the highest frequency.

```
summary(dat)
```

```
## transactions as itemMatrix in sparse format with
## 31101 rows (elements/itemsets/transactions) and
## 12767 columns (items) and a density of 0.0009590938
```

```
##
```

```
## most frequent items:
```

```
##   lord   thou   god   said   thy (Other)
```

```
## 6667   3881   3875   3602   3044 359755
```

```
##
```

```
## element (itemset/transaction) length distribution:
```

```
## sizes
```

```
##   2    3    4    5    6    7    8    9   10   11   12   13   14   15   16   17
```

```
## 13  235  550  840 1554 2258 2536 2611 2428 2465 2283 2139 1925 1751 1490 1248
```

```
## 18  19   20   21   22   23   24   25   26   27   28   29   30   31   32   33
```

```
## 1084 876 666 574 412 321 258 182 129 107 57 47 22 14 2 8
```

```
## 34  35  36  37
```

```
##      5      4      5      2
##
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##      2.00   8.00   12.00   12.24   15.00   37.00
##
## includes extended item information - examples:
##      labels
## 1      aaron
## 2      aaron's
## 3 aaronites
```

Remarks

- Itemset/transaction with size 9 has the highest frequency of 2611
- Itemset/transaction with the highest size 37 and size 32 have the lowest frequency of 2.

```
#Association Rule Analysis
rules <- apriori(dat, parameter = list(support = 0.01, confidence = 0.5,
    target = "rules", maxlen=5))

## Apriori
##
## Parameter specification:
## confidence minval smax arem aval originalSupport maxtime support minlen
##      0.5      0.1      1 none FALSE              TRUE        5      0.01      1
## maxlen target  ext
##      5  rules TRUE
##
## Algorithmic control:
## filter tree heap memopt load sort verbose
##    0.1 TRUE TRUE  FALSE TRUE    2    TRUE
##
## Absolute minimum support count: 311
##
## set item appearances ...[0 item(s)] done [0.00s].
## set transactions ...[12767 item(s), 31101 transaction(s)] done [0.17s].
## sorting and recoding items ... [230 item(s)] done [0.01s].
## creating transaction tree ... done [0.02s].
## checking subsets of size 1 2 3 4 done [0.01s].
## writing ... [29 rule(s)] done [0.00s].
## creating S4 object ... done [0.02s].
```

```
inspect(rules[1:5])
```

```
##      lhs      rhs      support      confidence coverage      lift      count
## [1] {answered} => {said} 0.01067490 0.6775510 0.01575512 5.850226 332
## [2] {art}      => {thou} 0.01434037 0.9867257 0.01453329 7.907280 446
## [3] {word}     => {lord} 0.01189672 0.5497771 0.02163918 2.564664 370
## [4] {moses}    => {lord} 0.01485483 0.5976714 0.02485451 2.788087 462
## [5] {she}      => {her}  0.01408315 0.6016484 0.02340761 15.684715 438
```

```
Rules <- as(rules, "data.frame")
head(Rules); tail(Rules)
```

```
##      rules      support confidence      coverage      lift count
## 1 {answered} => {said} 0.01067490 0.6775510 0.01575512 5.850226 332
## 2 {art}      => {thou} 0.01434037 0.9867257 0.01453329 7.907280 446
## 3 {word}     => {lord} 0.01189672 0.5497771 0.02163918 2.564664 370
## 4 {moses}    => {lord} 0.01485483 0.5976714 0.02485451 2.788087 462
## 5 {she}      => {her}  0.01408315 0.6016484 0.02340761 15.684715 438
## 6 {thus}     => {saith} 0.01462975 0.6435644 0.02273239 16.721383 455
```

```
##      rules      support confidence      coverage      lift count
## 24 {god,thee} => {lord} 0.01054628 0.5996344 0.01758786 2.797244 328
## 25 {god,thy}  => {thou} 0.01135012 0.5912898 0.01919552 4.738393 353
## 26 {god,thou} => {thy}  0.01135012 0.5064562 0.02241085 5.174539 353
## 27 {god,thy}  => {lord} 0.01340793 0.6984925 0.01919552 3.258409 417
## 28 {lord,thy} => {thou} 0.01530497 0.5157096 0.02967750 4.132720 476
## 29 {god,thou} => {lord} 0.01305424 0.5824964 0.02241085 2.717297 406
```

```
dim(Rules)
```

```
## [1] 29 6
```

```
inspect(rules[1:10], ruleSep = "---->", itemSep = " + ", setStart = "",
        setEnd = "",linebreak = FALSE)
```

```
##      lhs      rhs      support      confidence coverage      lift      count
## [1] answered ----> said 0.01067490 0.6775510 0.01575512 5.850226 332
## [2] art      ----> thou 0.01434037 0.9867257 0.01453329 7.907280 446
## [3] word     ----> lord 0.01189672 0.5497771 0.02163918 2.564664 370
## [4] moses    ----> lord 0.01485483 0.5976714 0.02485451 2.788087 462
## [5] she      ----> her  0.01408315 0.6016484 0.02340761 15.684715 438
```

```
## [6]  thus      ----> saith 0.01462975 0.6435644 0.02273239 16.721383 455
## [7]  thus      ----> lord 0.01626957 0.7157001 0.02273239 3.338682 506
## [8]  pass      ----> came 0.01488698 0.5758706 0.02585126 9.337932 463
## [9]  thine     ----> thy  0.01318286 0.5012225 0.02630141 5.121065 410
## [10] thine     ----> thou 0.01395454 0.5305623 0.02630141 4.251744 434
```

```
quality(rules[1:15])
```

##	support	confidence	coverage	lift	count
## 1	0.01067490	0.6775510	0.01575512	5.850226	332
## 2	0.01434037	0.9867257	0.01453329	7.907280	446
## 3	0.01189672	0.5497771	0.02163918	2.564664	370
## 4	0.01485483	0.5976714	0.02485451	2.788087	462
## 5	0.01408315	0.6016484	0.02340761	15.684715	438
## 6	0.01462975	0.6435644	0.02273239	16.721383	455
## 7	0.01626957	0.7157001	0.02273239	3.338682	506
## 8	0.01488698	0.5758706	0.02585126	9.337932	463
## 9	0.01318286	0.5012225	0.02630141	5.121065	410
## 10	0.01395454	0.5305623	0.02630141	4.251744	434
## 11	0.02684801	0.9835100	0.02729816	7.881511	835
## 12	0.01691264	0.5394872	0.03134947	2.516663	526
## 13	0.02819845	0.7326650	0.03848751	3.417821	877
## 14	0.03900196	0.9991763	0.03903411	8.007055	1213
## 15	0.02286100	0.5511628	0.04147777	6.012527	711

```
summary(rules)
```

```
## set of 29 rules
##
## rule length distribution (lhs + rhs):sizes
## 2 3
## 15 14
##
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##      2.000  2.000   2.000   2.483   3.000   3.000
##
## summary of quality measures:
##      support      confidence      coverage      lift
##      Min.    :0.01055    Min.    :0.5012    Min.    :0.01132    Min.    : 2.517
##      1st Qu.:0.01190    1st Qu.:0.5759    1st Qu.:0.01550    1st Qu.: 3.339
##      Median :0.01389    Median :0.6776    Median :0.02241    Median : 5.121
##      Mean   :0.01547    Mean   :0.7261    Mean   :0.02215    Mean   : 6.556
##      3rd Qu.:0.01514    3rd Qu.:0.9495    3rd Qu.:0.02630    3rd Qu.: 7.993
```

```
## Max. :0.03900 Max. :1.0000 Max. :0.04148 Max. :22.183
## count
## Min. : 328
## 1st Qu.: 370
## Median : 432
## Mean : 481
## 3rd Qu.: 471
## Max. :1213
##
## mining info:
## data ntransactions support confidence
## dat 31101 0.01 0.5
##
## apriori(data = dat, parameter = list(support = 0.01, confidence = 0.5, target = "rul
```

Remarks

- The parameters used for the R function `arules` are : `support = 0.01, confidence = 0.5, target = "rules", maxlen=5`.
- The maximum support is 0.03900 and the minimum support is 0.01055.
- The maximum confidence is 1.0000 and the minimum confidence is 0.5012.
- The maximum lift is 22.183 and the minimum lift is 2.517.

5.3 Top 5 rules in decreasing order of confidence (conf) for item sets of size/length 2 or 3.

```
rules0 <- data.frame(matrix(unlist(strsplit(as.character(Rules$rules), split="=>")),
  ncol=2, byrow=TRUE))
colnames(rules0) <- c("LHS", "RHS")
rule.size <- function(x){length(unlist(strsplit(as.character(x), split=",")))}
rules0$size <- apply(rules0, 1, rule.size)
```

```
z <- data.frame(Rules, size=rules0$size)
top.support <- z[order(z$confidence, decreasing = T),]
head(top.support, 5)
```

```
## rules support confidence coverage lift count size
## 20 {shalt,thee} => {thou} 0.01270056 1.0000000 0.01270056 8.013656 395 3
```

## 21	{shalt,thy} => {thou}	0.01514421	1.0000000	0.01514421	8.013656	471	3
## 14	{shalt} => {thou}	0.03900196	0.9991763	0.03903411	8.007055	1213	2
## 22	{lord,shalt} => {thou}	0.01225041	0.9973822	0.01228256	7.992678	381	3
## 2	{art} => {thou}	0.01434037	0.9867257	0.01453329	7.907280	446	2

Remarks

We observed that the rule (shalt) => (thou) is a creditable rule since it has a large level of support(0.03900196),large confidence(0.9991763) factor and a value of lift(8.007055) greater than 1. Thus, we expect to see **shalt** followed by **thou** in the King James Bible1. In other words, shalt and thou are words that commonly occur together in sentences.

5.4 Top 5 rules in decreasing order of the lift measure for item sets of size 2 or 3.

```
z <- data.frame(Rules, size=rules0$size)
top5.lift <- z[order(z$lift, decreasing = T),]
head(top5.lift, 5)
```

##	rules	support	confidence	coverage	lift	count	size
## 17	{lord,thus} => {saith}	0.01389023	0.8537549	0.01626957	22.182650	432	3
## 6	{thus} => {saith}	0.01462975	0.6435644	0.02273239	16.721383	455	2
## 5	{she} => {her}	0.01408315	0.6016484	0.02340761	15.684715	438	2
## 8	{pass} => {came}	0.01488698	0.5758706	0.02585126	9.337932	463	2
## 20	{shalt,thee} => {thou}	0.01270056	1.0000000	0.01270056	8.013656	395	3

Remarks

We observed that the rule (thus) => (saith) is a fairly a creditable rule since it has a fairly large level of support(0.01462975),fairly large confidence(0.6435644) factor and a value of lift(16.721383) greater than 1. Hence, thus and saith are words that commonly occur together in sentences in the King James Bible.

5.5 Conviction measures for the top-lift 5 rules in Part (d)

```
top5_liftrules <- sort(rules, decreasing = T, by='lift')[1:5,] # top lift 5 rules
interestMeasure(top5_liftrules, "conviction", transactions = dat)
```

```
## [1] 6.574666 2.697577 2.414051 2.212367      Inf
```

Remarks

- The problem associated with both the confidence and the lift measures is that they are not sensitive to rule direction. On the other hand, conviction is sensitive to rule direction. It attempts to measure the degree of implication of a rule. That is unlike lift, $\text{conviction}(A \Rightarrow B) \neq \text{conviction}(B \Rightarrow A)$.
- The $\text{conviction}((\text{shalt}, \text{thee}) \Rightarrow (\text{thou})) = \infty$. This is because the confidence obtained for the rule $(\text{shalt}, \text{thee}) \Rightarrow (\text{thou})$ in part b is 1.