

Project 2: Optimization and Kernel Trick

Appiah Prince* University of Texas at El Paso (UTEP)

September 27, 2022

Contents

1	Bring in the data	2
2	Exploratory Data Analysis (EDA)	3
2.1	Distinct levels or values for each variable	3
2.2	Missing Values	4
2.3	Parallel Boxplot of the Data	4
2.4	Bar plot of the binary response Class	5
3	Data Partitioning	6
4	Logistic Regression - Optimization	7
4.1	Pool the training data and the validation data together	7
4.2	Negative Likelihood Function and Test On the shill_bidding data	8
4.3	Standard error from the Hessian matrix	8
4.4	Convergence of the algorithm	8
4.5	Testing the significance of each attribute and table of results	9
4.6	Comparing results in 4(a) with fitting results from glm()	9
4.7	Making prediction using the test data	11

*pappiah@miners.utep.edu

5	Primitive LDA (The Kernel Trick)	11
5.1	Matrix of all predictors and Scaling X1 and X2	11
5.2	Train the primitive LDA classifier with D1 and use the prediction accuracy on D2	12
5.3	Apply the trained classifier with the ‘best’ kernel found in 5(b) to the test data D3.	13
5.4	Comparison of the prediction accuracy obtained in 4(c) and 5(c)	14

1 Bring in the data

```
data <- read.csv('Shill Bidding Dataset.csv')
names(data)
```

```
## [1] "Record_ID"          "Auction_ID"          "Bidder_ID"
## [4] "Bidder_Tendency"    "Bidding_Ratio"       "Successive_Outbidding"
## [7] "Last_Bidding"       "Auction_Bids"        "Starting_Price_Average"
## [10] "Early_Bidding"      "Winning_Ratio"       "Auction_Duration"
## [13] "Class"
```

```
# Remove the first three columns
shill_bidding <- data[, -c(1:3)]
names(shill_bidding)
```

```
## [1] "Bidder_Tendency"    "Bidding_Ratio"       "Successive_Outbidding"
## [4] "Last_Bidding"       "Auction_Bids"        "Starting_Price_Average"
## [7] "Early_Bidding"      "Winning_Ratio"       "Auction_Duration"
## [10] "Class"
```

```
#change the 0 value of the class variable to -1
shill_bidding$Class[shill_bidding$Class == 0] <- -1
#shill_bidding$Class <- ifelse(shill_bidding$Class == 0, -1, 1)
table(shill_bidding$Class)
```

```
##
## -1    1
## 5646  675
```

```
dim(shill_bidding)
```

```
## [1] 6321 10
```

Remarks

- The data Shill Bidding was loaded and the first three columns were removed since these were ID variables.
- The original dimension of the data was 6321 rows and 13 columns. However, since we removed three columns, our new dimension is 6,321 rows and 10 columns.
- I printed a table of the target variable(class) to confirm that the level 0 has been changed to -1.

2 Exploratory Data Analysis (EDA)

2.1 Distinct levels or values for each variable

```
aggregate(values ~ ind, unique(stack(shill_bidding)), length)
```

```
##               ind values
## 1      Bidder_Tendency  489
## 2      Bidding_Ratio   400
## 3 Successive_Outbidding    3
## 4      Last_Bidding  5807
## 5      Auction_Bids    49
## 6 Starting_Price_Average   22
## 7      Early_Bidding  5690
## 8      Winning_Ratio    72
## 9      Auction_Duration    5
## 10             Class     2
```

```
str(shill_bidding)
```

```
## 'data.frame':    6321 obs. of  10 variables:
## $ Bidder_Tendency      : num  0.2 0.0244 0.1429 0.1 0.0513 ...
## $ Bidding_Ratio        : num  0.4 0.2 0.2 0.2 0.222 ...
## $ Successive_Outbidding : num  0 0 0 0 0 0 0 1 1 0.5 ...
## $ Last_Bidding         : num  2.78e-05 1.31e-02 3.04e-03 9.75e-02 1.32e-03 ...
```

```
## $ Auction_Bids          : num  0 0 0 0 0 ...
## $ Starting_Price_Average: num  0.994 0.994 0.994 0.994 0 ...
## $ Early_Bidding         : num  2.78e-05 1.31e-02 3.04e-03 9.75e-02 1.24e-03 ...
## $ Winning_Ratio         : num  0.667 0.944 1 1 0.5 ...
## $ Auction_Duration      : int   5 5 5 5 7 7 7 7 7 7 ...
## $ Class                 : num  -1 -1 -1 -1 -1 -1 -1 1 1 1 ...
```

Remarks

The numerical variables **class** and **Successive_Outbidding** have only few distinct values.

2.2 Missing Values

```
library(questionr)
freq.na(shill_bidding)
```

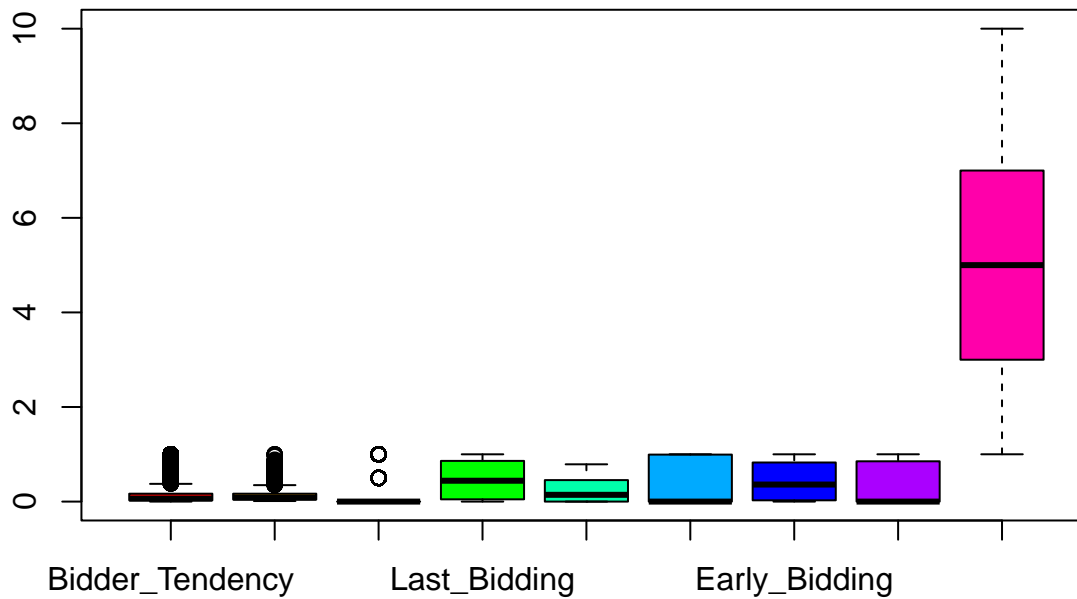
```
##                missing %
## Bidder_Tendency      0 0
## Bidding_Ratio        0 0
## Successive_Outbidding 0 0
## Last_Bidding         0 0
## Auction_Bids         0 0
## Starting_Price_Average 0 0
## Early_Bidding        0 0
## Winning_Ratio        0 0
## Auction_Duration     0 0
## Class                0 0
```

Remarks

There are no missing values in the data

2.3 Parallel Boxplot of the Data

```
boxplot(shill_bidding[, -10], col = rainbow(ncol(shill_bidding[, -10])))
```



Remarks

- The predictors have unequal range and unequal variation. In particular, the predictors **Auction_Duration**, **Starting_Price_Average**, **Winning_Ratio** and **Successive_Outbidding** have notable unequal range and variation.
- Hence, scaling is necessary for some modeling approaches.

2.4 Bar plot of the binary response Class


```
library(ggplot2)
library(dplyr)
```

```
##
## Attaching package: 'dplyr'

## The following objects are masked from 'package:stats':
##
##   filter, lag

## The following objects are masked from 'package:base':
##
##   intersect, setdiff, setequal, union
```

```
shill_bidding %>%  
  count(class = factor(Class)) %>%  
  mutate(pct = prop.table(n)) %>%  
  ggplot(aes(x = class, y = pct, label = scales::percent(pct), fill=class)) +  
  geom_col(position = 'dodge') +  
  geom_text(position = position_dodge(width = .9),  
            vjust = -0.5,  
            size = 3) +  
  scale_y_continuous(labels = scales::percent) +  
  theme(legend.position = "none")
```



Remarks

We see from the barplot that the percentage of 1's is 11% and percentage of -1's is 89%. Therefore, we do not have an unbalanced classification problem.

3 Data Partitioning

```
set.seed(125)
Data <- sample(seq(1, 3), size = nrow(shill_bidding), replace = TRUE,
               prob = c(0.5, 0.25, 0.25))
train_data <- shill_bidding[Data == 1, ] # training data
validation_data <- shill_bidding[Data == 2, ] # validation data
test_data <- shill_bidding[Data == 3, ] # test data

dim(train_data)
```

```
## [1] 3171  10
```

```
dim(validation_data)
```

```
## [1] 1596  10
```

```
dim(test_data)
```

```
## [1] 1554  10
```

Remarks

- There are 3171 observations and 10 variables in the training data.
- There are 1596 observations and 10 variables in the validation data.
- There are 1554 observations and 10 variables in the test data.

4 Logistic Regression - Optimization

- Part 4(a)

4.1 Pool the training data and the validation data together

```
train_valid_data <- rbind(train_data, validation_data)
#head(train_valid_data)
dim(train_valid_data)
```

```
## [1] 4767  10
```

4.2 Negative Likelihood Function and Test On the shill_bidding data

```
# THE NEGATIVE LOGLIKELIHOOD FUNCTION FOR Y=+1/-1
nloglik <- function(beta, X, y){
  if (length(unique(y)) !=2) stop("Are you sure you've got Binary Target?")
  X <- cbind(1, X)
  nloglik <- sum(log(1+ exp(-y*X%%beta)))
  return(nloglik)
}

y <- train_valid_data$Class
X <- as.matrix(train_valid_data[, c(1:9)])
p <- NCOL(X) +1
fit <- optim(par=rep(0,p), fn=nloglik, method="BFGS", X=X, y=y,
            hessian = TRUE)
beta.hat <- fit$par # obtaining the regression parameters
beta.hat
```

```
## [1] -10.10785135  1.05384204  1.25123047  10.49377567  0.93247462
## [6]  0.63452561  0.12535724 -0.64309290  4.77652987  0.05764265
```

Remarks

The optimization method that was employed in R function `optim()` is BFGS

4.3 Standard error from the Hessian matrix

```
hessian <- fit$hessian # Hessian matrix
inv_hessian <- solve(hessian)
standard_error <- sqrt(diag(inv_hessian))
standard_error
```

```
## [1] 0.77083290 0.53101537 0.96953407 0.64970474 0.77500494 0.73158830
## [7] 0.33141911 0.77395553 0.63178907 0.05016893
```

4.4 Convergence of the algorithm


```
fit$convergence
```

```
## [1] 0
```

Remarks The algorithm converges since the output of `fit$convergence` is 0.

4.5 Testing the significance of each attribute and table of results

```
p0 <- length(beta.hat)-1
z.wald <- beta.hat/standard_error
pvalue <- pchisq(z.wald^2, df=1, lower.tail=FALSE)
result <- data.frame(beta.hat, standard_error, z.wald, pvalue)
row.names(result) <- c("Intercept", names(shill_bidding[, -10]))
round(result, digits = 4)
```

##	beta.hat	standard_error	z.wald	pvalue
## Intercept	-10.1079	0.7708	-13.1129	0.0000
## Bidder_Tendency	1.0538	0.5310	1.9846	0.0472
## Bidding_Ratio	1.2512	0.9695	1.2905	0.1969
## Successive_Outbidding	10.4938	0.6497	16.1516	0.0000
## Last_Bidding	0.9325	0.7750	1.2032	0.2289
## Auction_Bids	0.6345	0.7316	0.8673	0.3858
## Starting_Price_Average	0.1254	0.3314	0.3782	0.7052
## Early_Bidding	-0.6431	0.7740	-0.8309	0.4060
## Winning_Ratio	4.7765	0.6318	7.5603	0.0000
## Auction_Duration	0.0576	0.0502	1.1490	0.2506

Remarks

- Taken $\alpha = 0.05$ as a threshold, we see from our results that the p-values for the predictors **Bidder_Tendency**, **Successive_Outbidding** and **Winning_Ratio** are less than $\alpha = 0.05$. Hence, these attributes are statistically significant.
- Part (4b)

4.6 Comparing results in 4(a) with fitting results from `glm()`

```

y <- factor(train_valid_data$Class)
fit.logit <- glm(y~Bidder_Tendency+Bidding_Ratio+Successive_Outbidding+
                Last_Bidding+Auction_Bids+Starting_Price_Average+
                Early_Bidding+Winning_Ratio+Auction_Duration,
                data=train_valid_data, family=binomial(link = "logit"))
result <- summary(fit.logit)
round(result$coefficients, 4)

```

##	Estimate	Std. Error	z value	Pr(> z)
## (Intercept)	-10.1079	0.7708	-13.1135	0.0000
## Bidder_Tendency	1.0539	0.5310	1.9847	0.0472
## Bidding_Ratio	1.2512	0.9695	1.2906	0.1969
## Successive_Outbidding	10.4938	0.6497	16.1522	0.0000
## Last_Bidding	0.9325	0.7750	1.2032	0.2289
## Auction_Bids	0.6345	0.7316	0.8673	0.3858
## Starting_Price_Average	0.1254	0.3314	0.3782	0.7052
## Early_Bidding	-0.6431	0.7739	-0.8309	0.4060
## Winning_Ratio	4.7765	0.6318	7.5606	0.0000
## Auction_Duration	0.0576	0.0502	1.1490	0.2506

```
fit.logit$converged
```

```
## [1] TRUE
```

Remarks

- The `glm()` result also converges.
- Taken $\alpha = 0.05$ as a threshold, we see from our results that the p-values for the attributes **Bidder_Tendency**, **Successive_Outbidding** and **Winning_Ratio** are less than $\alpha = 0.05$. Hence, these attributes are statistically significant.
- The coefficient of **Winning_Ratio** says that, holding the other predictors at a fixed value, we will see $e^{4.7765} = 11868.82\%$ increase in the odds of getting into a positive Class for a unit increase in **Winning_Ratio**.
- There appears to be no difference between the results of the two methods.
- Part 4(c)

4.7 Making prediction using the test data

```
my_fun <- function(x){
  exp(x)/(1+exp(x) )
}

test_data <- test_data
new_X <- as.matrix(cbind(1,test_data[, -10]))
y_hat_prime <- sign(my_fun(new_X%%fit$par) - 0.5)
conf_matt <- table(test_data$Class, y_hat_prime) # gives the confusion matrix
#round(mean(test_data$Class == y_hat_prime),)
pred_acc <- sum(diag(conf_matt))/sum(conf_matt) # gives the accuracy
pred_acc

## [1] 0.9787645
```

Remarks

With a threshold of 0.5, our prediction accuracy is 0.9787645. This means, the algorithm in 4(a) predicts or models the data very well.

5 Primitive LDA (The Kernel Trick)

5.1 Matrix of all predictors and Scaling X1 and X2

- Part 5(a)

```
# matrix of all predictors for the three train, validation and test sets
X1 <- as.matrix(train_data[, -10])
X2 <- as.matrix(validation_data[, -10])
X3 <- as.matrix(test_data[, -10])

#scale X1
X1_scale <- scale(X1, center = TRUE, scale = TRUE)

# scale X2 according to the column means and SDs computed from X1
mean_X1 <- attributes(X1_scale)$'scaled:center'
sd_X1 <- attributes(X1_scale)$'scaled:scale'
X2_scale <- scale(X2, center = mean_X1, scale = sd_X1)
```

5.2 Train the primitive LDA classifier with D1 and use the prediction accuracy on D2

- Part 5(b)

```
library(kernlab)

##
## Attaching package: 'kernlab'


## The following object is masked from 'package:ggplot2':
##
##      alpha

LDA_P <- function (kernel, X, Y=NULL, target) {
  kernmat <- kernelMatrix
  w.z <- colMeans(kernmat(kernel, x=X[target==1,], y=Y)) -
    colMeans(kernmat(kernel, x=X[target==1,], y=Y))
  b <- (mean(kernmat(kernel, X[target==1,])) -
    mean(kernmat(kernel, X[target==1,])))*.5
  yhat <- sign(w.z + b)
  return(yhat)
}

deg_vec <- 1:15
pred.acc_vec <- rep(0, length(deg_vec))
for (i in 1:length(deg_vec)) {
  d <- deg_vec[i]
  kern <- polydot(degree = d, offset = 1, scale = 1)

  #compute prediction accuracy
  ypred <- LDA_P(kern, X1_scale, X2_scale, train_data$Class)
  yobserved <- validation_data$Class
  conf_mat <- table(ypred, yobserved)
  pred_accuracy <- sum(diag(conf_mat))/sum(conf_mat)
  pred.acc_vec[i] <- pred_accuracy
}

plot(deg_vec, pred.acc_vec, type = "b", col="red")
```



```
max(pred.acc_vec); min(pred.acc_vec)
```

```
## [1] 0.9655388
```

```
## [1] 0.8997494
```

Remarks

- The kernel family used was polynomial kernel.
- From the plot of the prediction accuracy values versus the candidate parameter values, the best choice of our parameter (degree) is 2.
- The maximum value of the prediction accuracy is **0.9655388**. Thus, we see that the polynomial kernel helps well in the classification.
- The minimum value of the prediction accuracy is **0.8997494**.

5.3 Apply the trained classifier with the ‘best’ kernel found in 5(b) to the test data D3.

- Part 5(c)

```

# Scale X.prime
X_prime <- as.matrix(train_valid_data[, -10])
X_prime_scale <- scale(X_prime, center = TRUE, scale = TRUE)

# Scale X3 according to the column means and SDs computed from X.prime
mean_X_prime_scale <- attributes(X_prime_scale)$'scaled:center'
sd_X_prime_scale <- attributes(X_prime_scale)$'scaled:scale'
X3_scale <- scale(X3, center = mean_X_prime_scale , scale = sd_X_prime_scale)

# Apply the best kernel to the test data
kern <- polydot(degree = which.max(pred.acc_vec))

#compute prediction accuracy
ypred <- LDA_P(kern, X_prime_scale, X3_scale, train_valid_data$Class)
yobserved <- test_data$Class
conf_mat <- table(ypred, yobserved)
pred_accuracy_test <- sum(diag(conf_mat))/sum(conf_mat)
pred_accuracy_test

```

```
## [1] 0.972973
```

Remarks

- The prediction accuracy obtained after applying the trained classifier with the ‘best’ kernel found in 5(b) to the test data is **0.972973**. This shows an excellent prediction or classification ability of the model.

5.4 Comparison of the prediction accuracy obtained in 4(c) and 5(c)

```

comparison <- data.frame(pred_accuracy_test, pred_acc )
knitr::kable(comparison,
              col.names = c("Primitive LDA", "Logistic regression (optimization)"),
              caption = "Prediction accuracy", align = "cc")

```

Table 1: Prediction accuracy

Primitive LDA	Logistic regression (optimization)
0.972973	0.9787645

Remarks

There is a small difference between the prediction accuracy obtained in 4(c) and 5(c)